

# Chapter 7 - Multi-step Bootstrapping

Atul Balaji - EE16B002

February 26, 2019

- MC methods use an exact sample of return (till episode termination) whereas TD methods use an estimate using a one-step look ahead. n-step methods generalize both methods.
- In TD updates are performed for each time step by bootstrapping on the estimate of the next-state's value. But, bootstrapping works better if the estimates differ over a large temporal range. This is because successive states might be correlated and as a result, there will be only a small difference between successive state's value functions. Therefore, the increments will be less in magnitude and learning will be slower.
- By contrast, n-step returns consider bootstrapping with a state that occurs n time steps in the future. Thus, there would be lesser correlation between the states  $S_t$  and  $S_{t+n}$ . This results in better learning.

## 1 n-step TD

- Given state s, the target for n-step backup is the n-step return  $G_{t:t+n}$ .

$$G_{t:t+n} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V_{t+n-1}, \text{ if } n \geq 1, 0 \leq t < T - n. \quad (1)$$

$$G_{t:t+n} = G_t \text{ if } t + n \geq T. \quad (2)$$

- The target depends on the rewards obtained over n-steps, so, no update will happen during the first (n-1) steps in the episode. To make up for this, the same number of updates are done once the episode terminates.
- As  $n \rightarrow \infty$ , the n-step TD becomes the Monte Carlo update. MC method converges to an estimate which minimizes error on the training data. This means n-step TD gets closer to MC as  $n \uparrow$  and hence error on training data for n-step TD is less than one-step TD. This is referred to as the error-reduction property which bounds the max-error in the estimate of the value function.

## 2 n-step SARSA

- Here, n-step methods are combined with SARSA to produce on-policy TD control. The target is approximated using n rewards and is bootstrapped with an estimate of  $Q(S_{t+n}, A_{t+n})$ .

$$G_{t:t+n} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n Q_{t+n-1}(S_{t+n}, A_{t+n}), \text{ } n \geq 1, 0 \leq t < T - n. \quad (3)$$

$$G_{t:t+n} = G_t \text{ if } t + n \geq T \quad (4)$$

- After n steps, update is done for  $Q(S_t, A_t)$ . For  $s \neq S_t$  and  $a \neq A_t$ ,  $Q(s, a)$  remain unchanged.
- For Expected SARSA, instead of using  $Q(S_{t+n}, A_{t+n})$ , we use expectation over  $Q(S_{t+n}, a)$  for all actions  $a$ .

## 3 n-step Off-policy Learning

- Say we are using a behaviour policy  $b$  which could be some  $\epsilon$ -greedy policy and we estimate the target policy  $\pi$ . This method uses importance sampling.
- An off policy version of n-step TD can be made by weighing the samples with the relative probability of taking that action according to policies  $\pi$  and  $b$ . The weights here are the **importance sampling ratios**  $\rho_{t:t+n-1}$ . Here, since n steps are considered, the weight depends on relative probability of those actions  $A_t$  to  $A_{t+n-1}$ .
- Similarly, off-policy versions of n-step SARSA, n-step expected SARSA can be derived.

## 4 n-step Tree Backup algorithm

- Consider a particular  $(S_t, A_t)$ . In this method, the target consists of a backup using  $Q(S_{t+1}, x)$  where  $x \neq A_{t+1}$  (not selected). It uses the sample  $(S_{t+1}, A_{t+1})$  to bootstrap in order to calculate the **n-step return**  $G_{t:t+n}$ . Each first level action  $a$  contributes with a weight of  $\pi(a|S_{t+1})$ . The probabilities get multiplied in the subsequent levels.
- A state-action tree is formed, with each node (actions not selected) contributing to the target, its probability of occurring under policy  $\pi$ .
- The n-step return can be found using a recursive formula:

$$G_{t:t+n} = R_{t+1} + \gamma \sum_{a \neq A_{t+1}} \pi(a|S_{t+1}) Q_{t+n-1}(S_{t+1}, a) + \gamma \pi(A_{t+1}|S_{t+1}) G_{t+1:t+n}, \quad (5)$$

for  $t < T - 1, n \geq 2$ .

- The action value update is done only for the state-action pair  $(S_t, A_t)$  and the equation can be written as:

$$Q_{t+n}(S_t, A_t) = Q_{t+n-1}(S_t, A_t) + \alpha [G_{t:t+n} - Q_{t+n-1}(S_t, A_t)] \text{ for } 0 \leq t < T \quad (6)$$

- The n-step tree backup estimate has lower variance compared to importance sampling based methods.